



Elementary Graph Algorithms

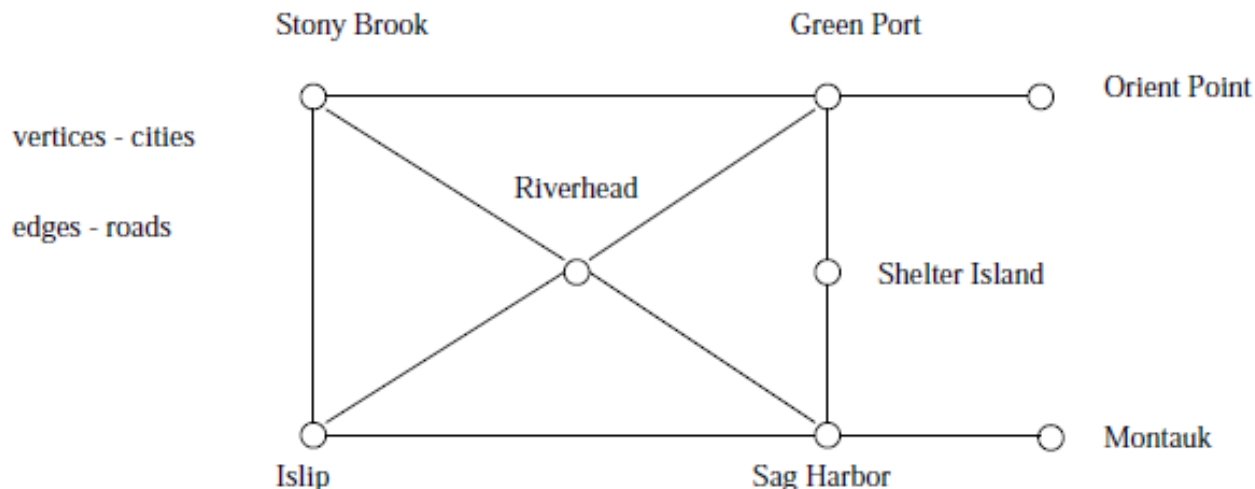
Graphs

- Graphs are one of the unifying themes of computer science. A graph $G = (V, E)$ is defined by a set of vertices V , and a set of edges E consisting of ordered or unordered pairs of vertices from V .
- Thus a graph $G = (V, E)$
 - V = set of vertices
 - E = set of edges = subset of $V \times V$
 - Thus $|E| = O(|V|^2)$

Examples

- **Road Networks**


- In modeling a road network, the vertices may represent the cities or junctions, certain pairs of which are connected by roads/edges.



Examples


- **Electronic Circuits**

- In an electronic circuit, with junctions as vertices and components as edges.



Graph Variations


- Variations:
 - A *connected graph* has a path from every vertex to every other
 - In an *undirected graph*:
 - Edge $(u,v) = \text{Edge}(v,u)$
 - No self-loops



Graph Variations

■ In a *directed* graph:


- Edge (u,v) goes from vertex u to vertex v , notated $u \rightarrow v$
- (u,v) is different from (v,u)




- Road networks *between cities* are typically undirected.
- Street networks *within* cities are almost always directed because of one-way streets.
- Most graphs of graph-theoretic interest are undirected.

Graph Variations

- More variations:
 - A *weighted graph* associates weights with either the edges or the vertices
 - E.g., a road map: edges might be weighted with their length, drive-time or speed limit.
 - In *unweighted* graphs, there is no cost distinction between various edges and vertices.



unweighted




weighted


-
- A *multigraph* allows multiple edges between the same vertices
 - E.g., the call graph in a program (a function can get called from multiple points in another function)

Simple vs. Non-simple Graphs

- Certain types of edges complicate the task of working with graphs. A *self-loop* is an edge $(x; x)$ involving only one vertex.
- An edge $(x; y)$ is a *multi-edge* if it occurs more than once in the graph.
- Any graph which avoids these structures is called *simple*.




simple




non-simple

Spars Vs Dense Graphs

- Graphs are *sparse* when only a small fraction of the possible number of vertex pairs actually have edges defined between them.
- Graphs are usually sparse due to application-specific constraints.
- Road networks must be sparse because of road junctions.



sparse




dense

Spars Vs Dense Graphs


- Typically dense graphs have a quadratic number of edges while sparse graphs are linear in size.

Cyclic vs. Acyclic Graphs

- An *acyclic* graph does not contain any cycles. *Trees* are connected acyclic *undirected* graphs.
- Directed acyclic graphs are called *DAGs*. They arise naturally in scheduling problems, where a directed edge $(x; y)$ indicates that x must occur before y





cyclic



acyclic


Labeled vs. Unlabeled Graphs

- In *labeled graphs*, each vertex is assigned a unique name or identifier to distinguish it from all other vertices.
- An important graph problem is *isomorphism testing*, determining whether the topological structure of two graphs are in fact identical if we ignore any labels.




Subgraphs

- A subgraph S of a graph G is a graph such that
 - The edges of S are a subset of the edges of G
- A spanning subgraph of G is a subgraph that contains all the vertices of G



Subgraph



Spanning subgraph


Trees and Forests

- A tree is an undirected graph T such that


- T is connected
- T has no cycles

This definition of tree is different from the one of a rooted tree

- A forest is an undirected graph without cycles
- The connected components of a forest are trees



Tree



Forest

Graph in Applications


- Internet: web graphs
 - Each page is a vertex
 - Each edge represent a hyperlink
 - Directed
- GPS: highway maps
 - Each city is a vertex
 - Each freeway segment is an undirected edge
 - Undirected
- Graphs are ubiquitous in computer science

Representing Graphs

- Assume $V = \{1, 2, \dots, n\}$
- An *adjacency matrix* represents the graph as a $n \times n$ matrix A :
 - $A[i, j] = 1$ if edge $(i, j) \in E$ (or weight of edge)
 $= 0$ if edge $(i, j) \notin E$

Graphs: Adjacency Matrix


- Example:



A	1	2	3	4
1				
2		?	?	
3				
4				

Graphs: Adjacency Matrix

- Example:




A	1	2	3	4
1	0	1	1	0
2	0	0	1	0
3	0	0	0	0
4	0	0	1	0

Graphs: Adjacency Matrix

- The adjacency matrix is a dense representation
 - Usually too much storage for large graphs
 - But can be very efficient for small graphs
- Most large interesting graphs are sparse
 - E.g., planar graphs, in which no edges cross
 - For this reason the *adjacency list* is often a more appropriate representation

Graphs: Adjacency List


- Adjacency list: for each vertex $v \in V$, store a list of vertices adjacent to v
- Example:
 - $\text{Adj}[1] = \{2, 3\}$
 - $\text{Adj}[2] = \{3\}$
 - $\text{Adj}[3] = \{\}$
 - $\text{Adj}[4] = \{3\}$
- Variation: can also keep a list of edges coming *into* vertex




Graphs: Adjacency List

- How much storage is required?
 - The *degree* of a vertex v in an undirected graph = # incident edges
 - Directed graphs have in-degree, out-degree
 - For directed graphs, # of items in adjacency lists is
$$\sum \text{out-degree}(v) = |E|$$
takes $\Theta(V + E)$ storage
 - For undirected graphs, # items in adj lists is
$$\sum \text{degree}(v) = 2 |E|$$
also $\Theta(V + E)$ storage
- So: Adjacency lists take $O(V+E)$ storage

Graph Representation



(a)




(b)


	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

(c)

Figure 22.1 Two representations of an undirected graph. (a) An undirected graph G having five vertices and seven edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .



(a)



(b)

	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

(c)

Figure 22.2 Two representations of a directed graph. (a) A directed graph G having six vertices and eight edges. (b) An adjacency-list representation of G . (c) The adjacency-matrix representation of G .

Google Problem: Graph Searching

- How to search and explore in the Web Space?
- How to find all web-pages and all the hyperlinks?
 - Start from one vertex “www.google.com”
 - Systematically follow hyperlinks from the discovered vertex/web page
 - Build a search tree along the exploration

General Graph Searching


- Given: a graph $G = (V, E)$, directed or undirected
- Goal: methodically explore every vertex and every edge
- Ultimately: build a tree on the graph
 - Pick a vertex as the root
 - Choose certain edges to produce a tree
 - Note: might also build a *forest* if graph is not connected

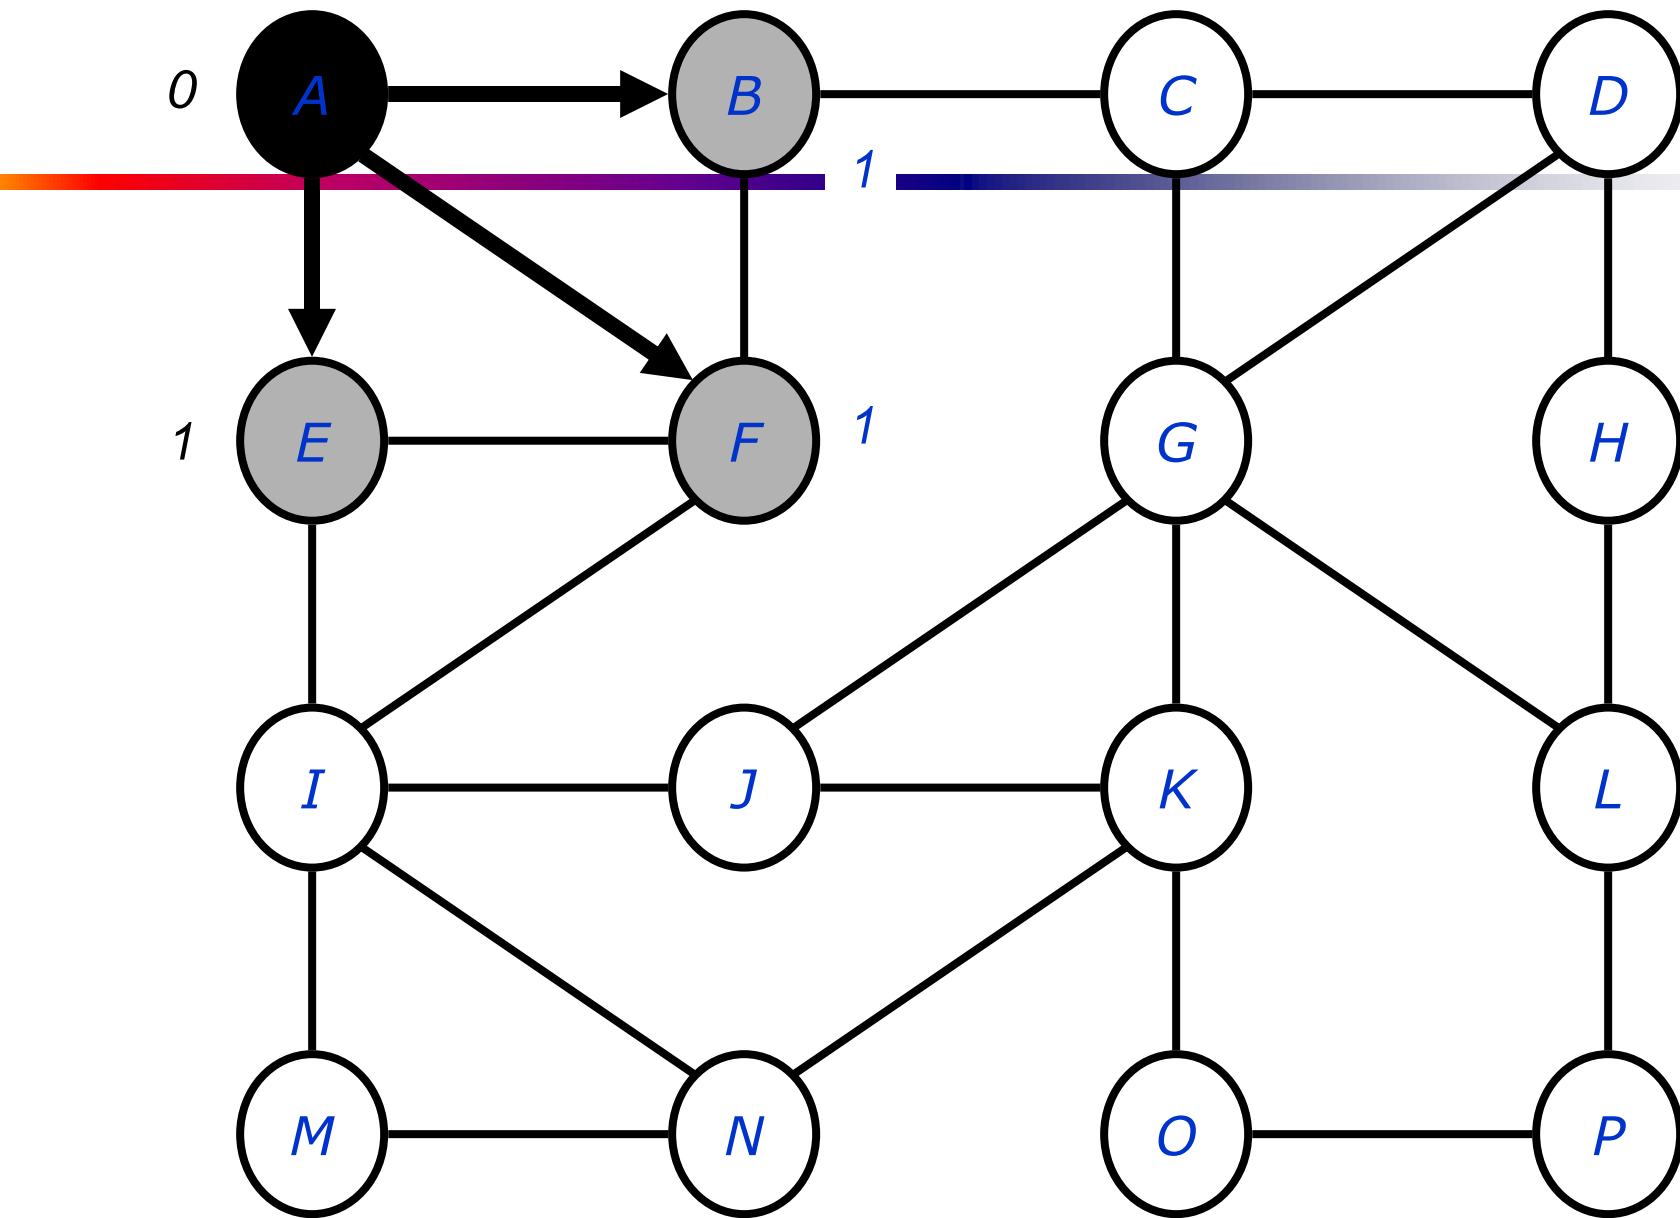
Breadth-First Search

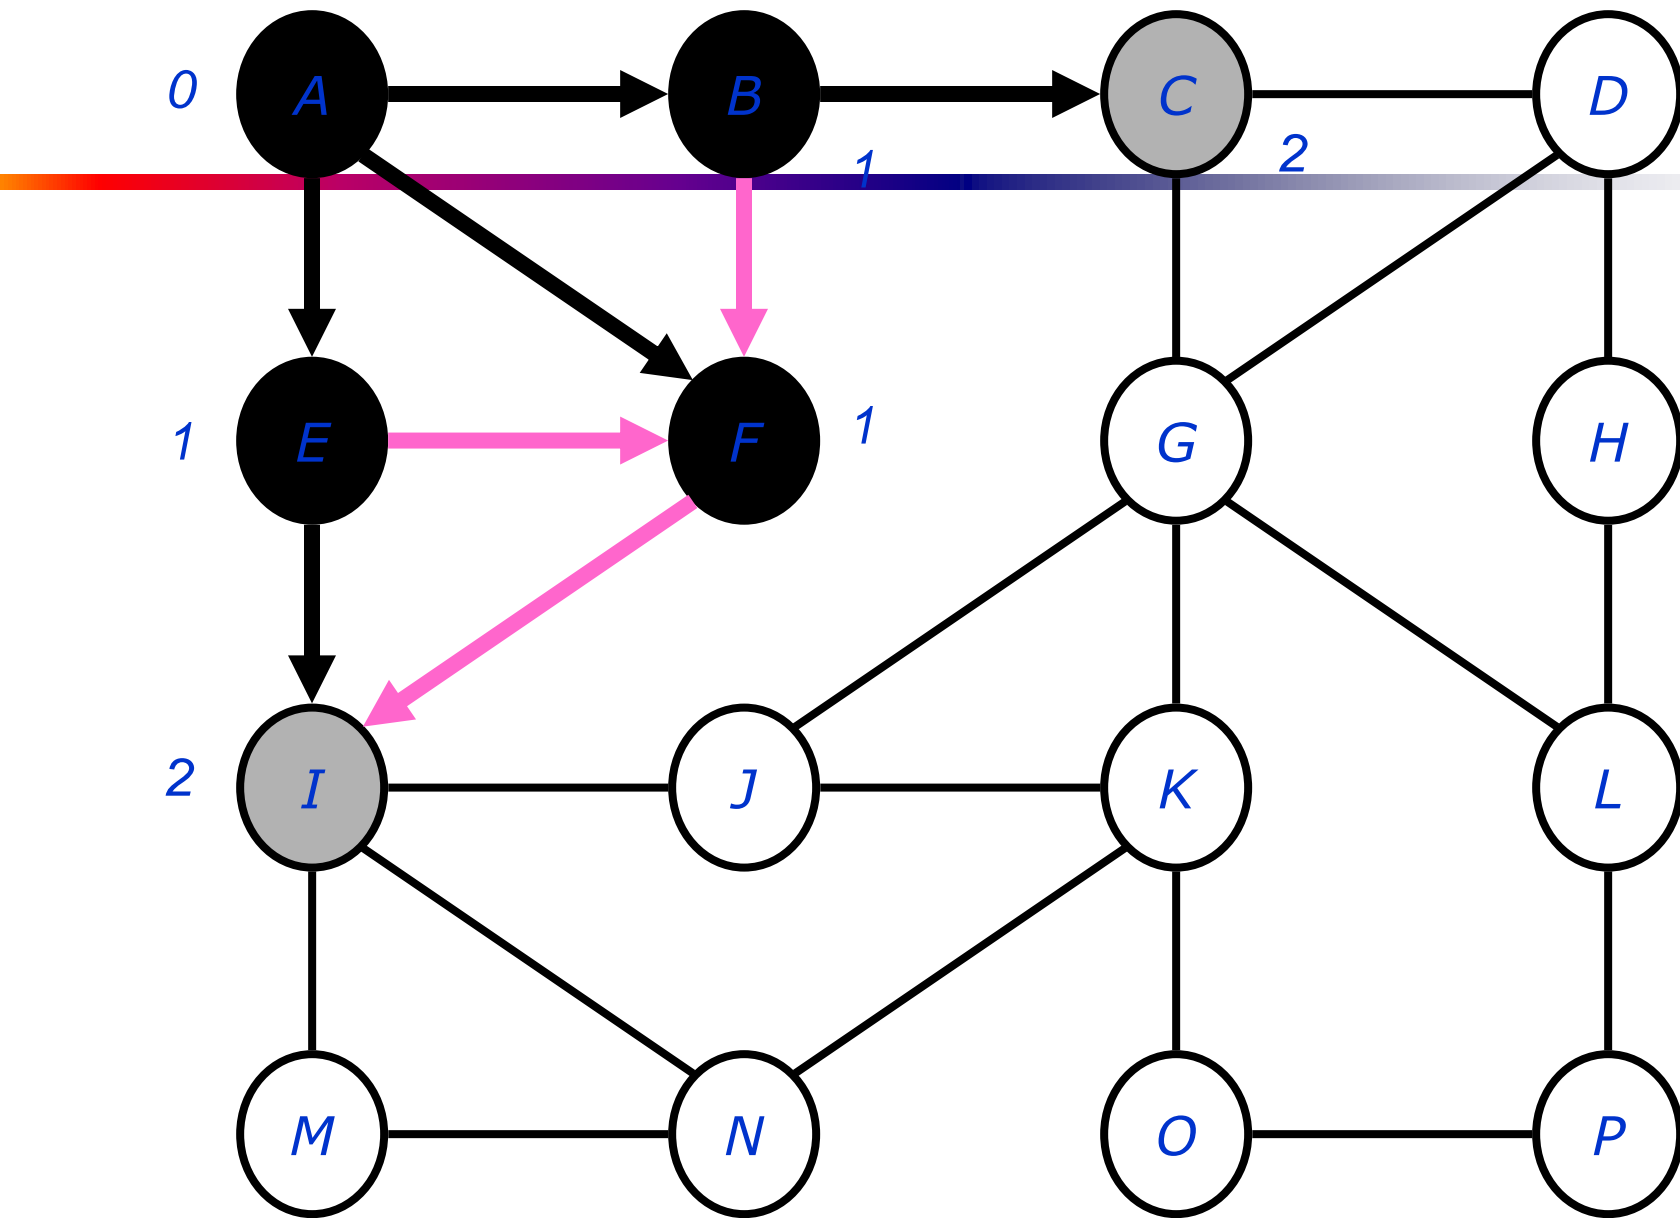
- “Explore” a graph, turning it into a tree
 - One vertex at a time
 - Expand frontier of explored vertices across the *breadth* of the frontier
- Builds a tree over the graph
 - Pick a *source vertex* to be the root
 - Find (“discover”) its children, then their children, etc.


Breadth-First Search

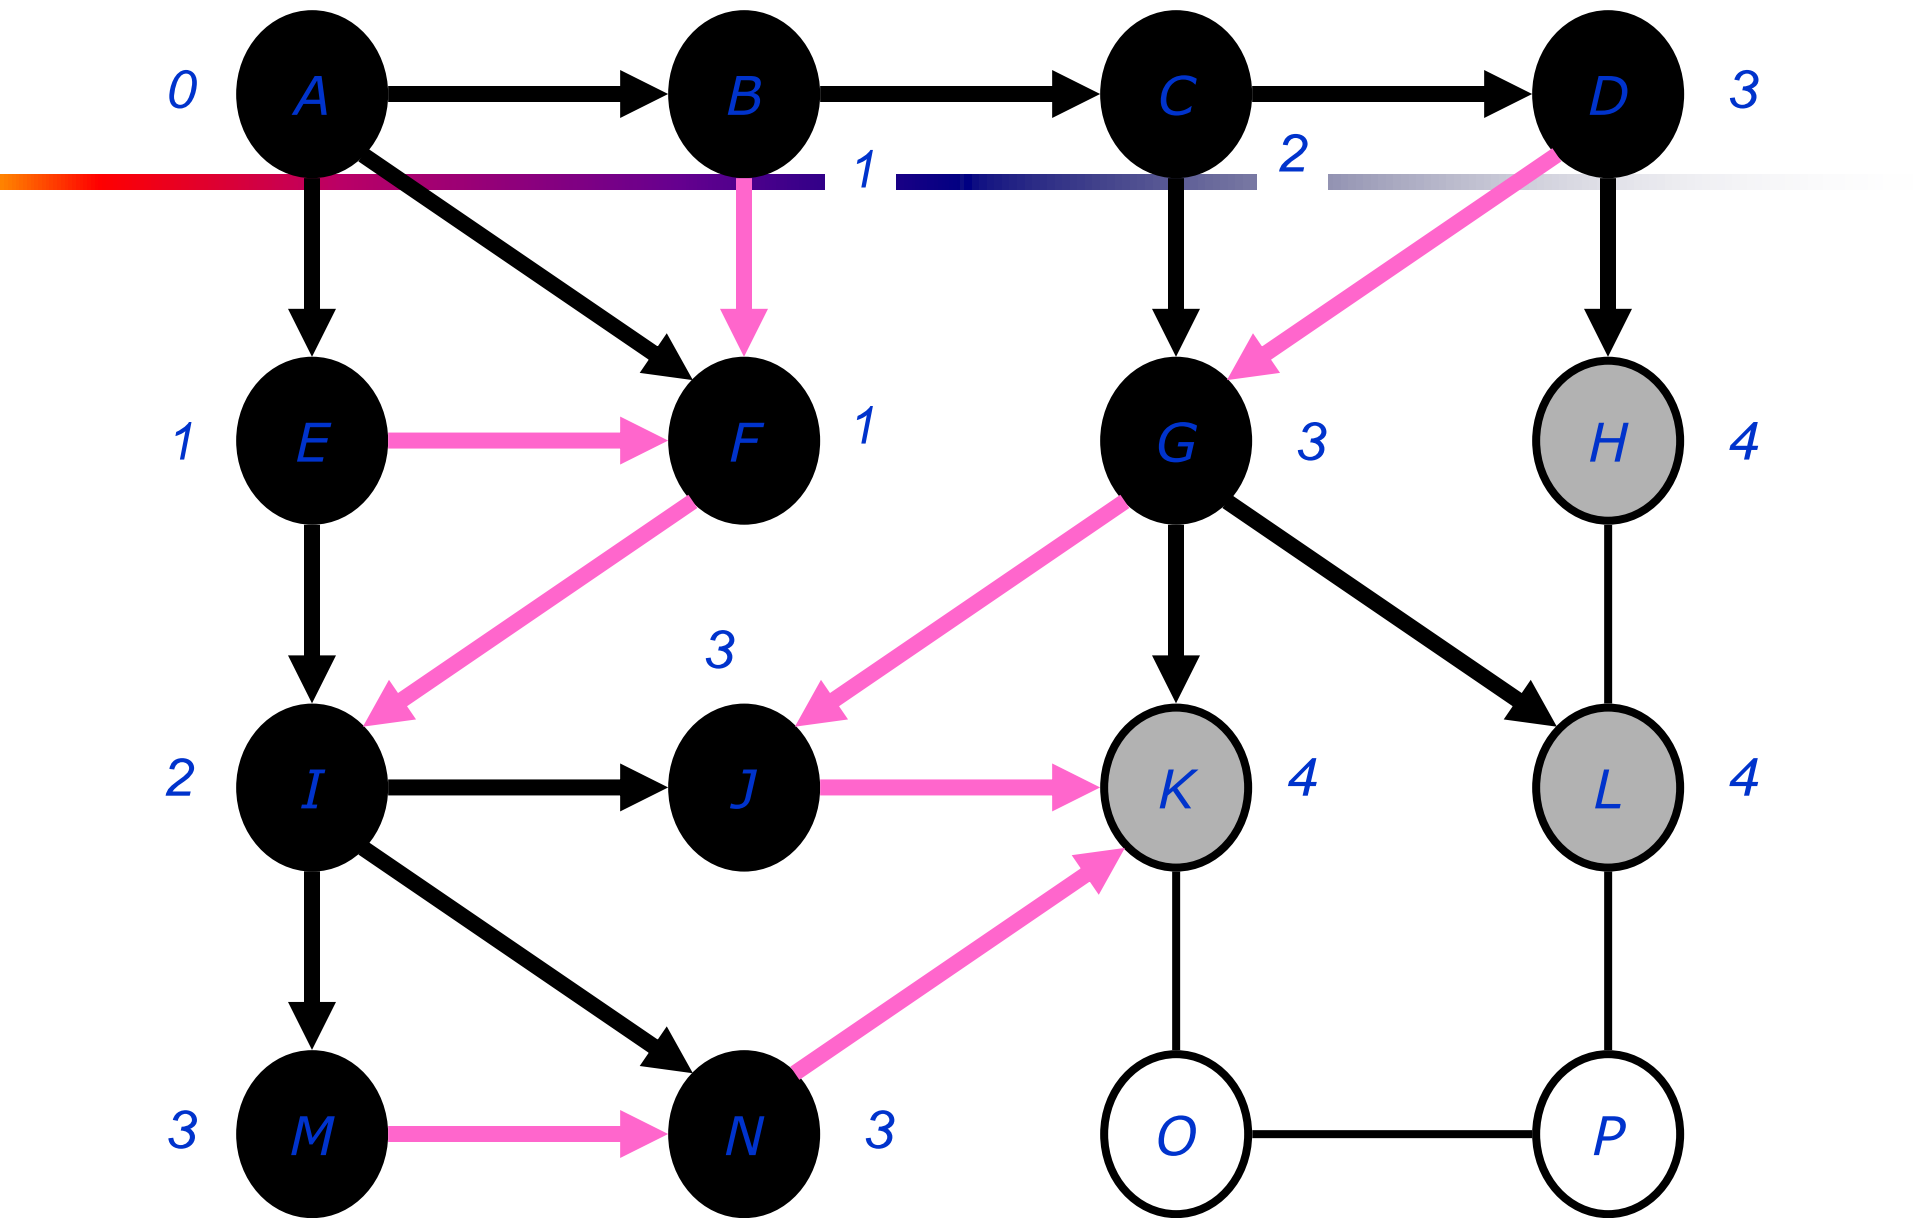
- Again will associate vertex “colors” to guide the algorithm
 - White vertices have not been discovered
 - All vertices start out white
 - Grey vertices are discovered but not fully explored
 - They may be adjacent to white vertices
 - Black vertices are discovered and fully explored
 - They are adjacent only to black and gray vertices
- Explore vertices by scanning adjacency list of grey vertices







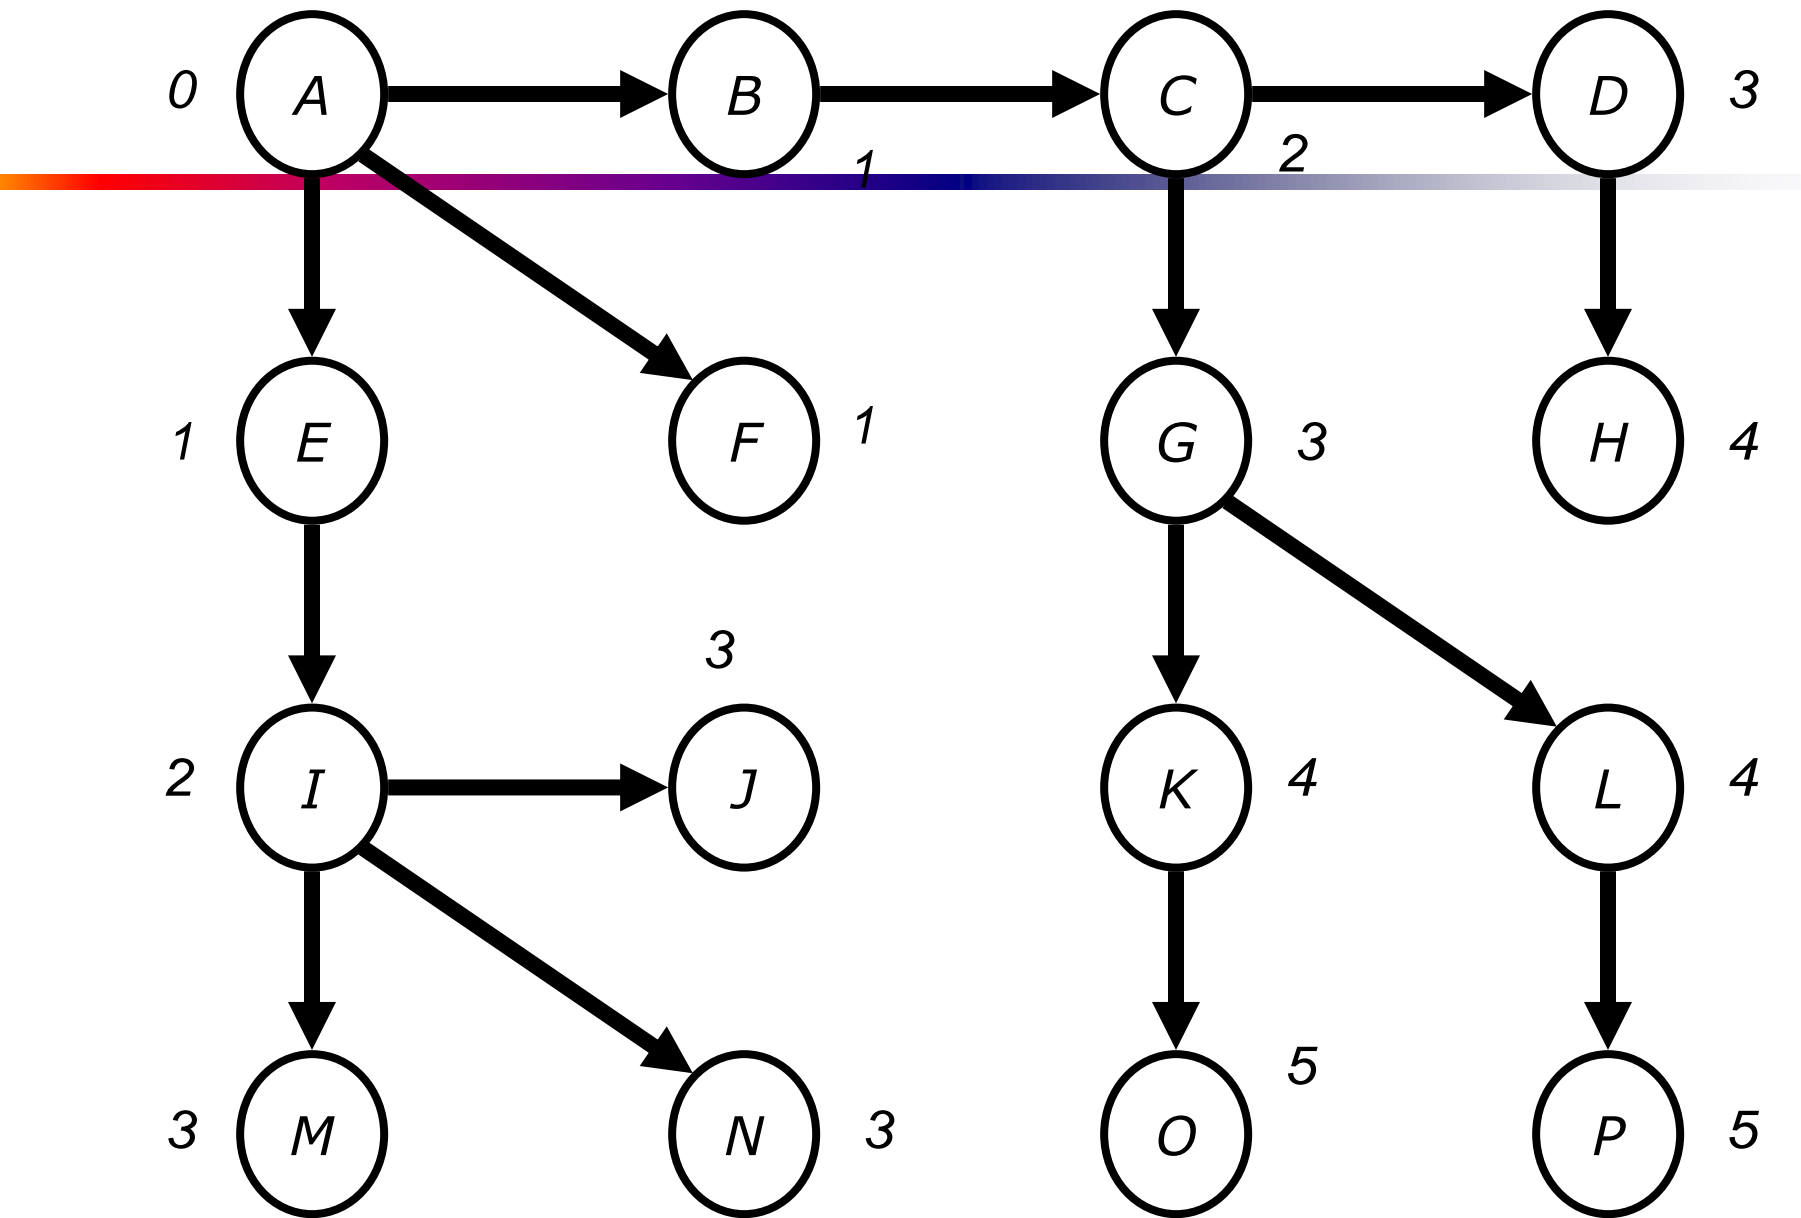














Breadth-First Search

```
BFS(G, s) {  
    initialize vertices;  
    Q = {s};           // Q is a queue (duh); initialize to s  
    while (Q not empty) {  
        u = RemoveTop(Q);  
        for each v ∈ u->adj {  
            if (v->color == WHITE)  
                v->color = GREY;  
                v->d = u->d + 1;  
                v->p = u;  
                Enqueue(Q, v);  
        }  
        u->color = BLACK;  
    }  
}
```

Breadth-First Search: Example




Breadth-First Search: Example



$Q:$ s


Breadth-First Search: Example



Q :

w	r
-----	-----


Breadth-First Search: Example



Q :

r	t	x
-----	-----	-----


Breadth-First Search: Example



Q :

t	x	v
-----	-----	-----


Breadth-First Search: Example



Q :

x	v	u
-----	-----	-----


Breadth-First Search: Example



Q :

v	u	y
-----	-----	-----


Breadth-First Search: Example



Q :


u	y
-----	-----

Breadth-First Search: Example






Q : y

Breadth-First Search: Example



$Q: \emptyset$

BFS: The Code Again

```
BFS(G, s) {  
    initialize vertices;  Touch every vertex:  $O(V)$   
    Q = {s};  
    while (Q not empty) {  
        u = RemoveTop(Q);  u = every vertex, but only once  
        for each v  $\in$  u->adj {  
             So v = every vertex that appears in some other vert's adjacency list  
            if (v->color == WHITE)  
                v->color = GREY;  
                v->d = u->d + 1;  
                v->p = u;  
                Enqueue(Q, v);  
            }  
        u->color = BLACK;  
    }  
}
```

What will be the running time?


Total running time: $O(V+E)$



Dijkstra's Algorithm

Single-Source Shortest Path Problem

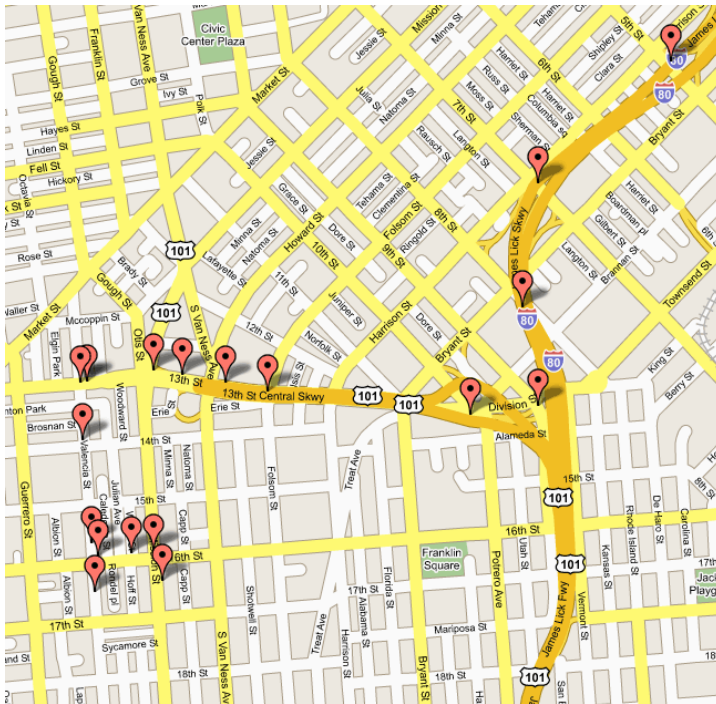
Single-Source Shortest Path Problem - The problem of finding shortest paths from a source vertex v to all other vertices in the graph.



Applications

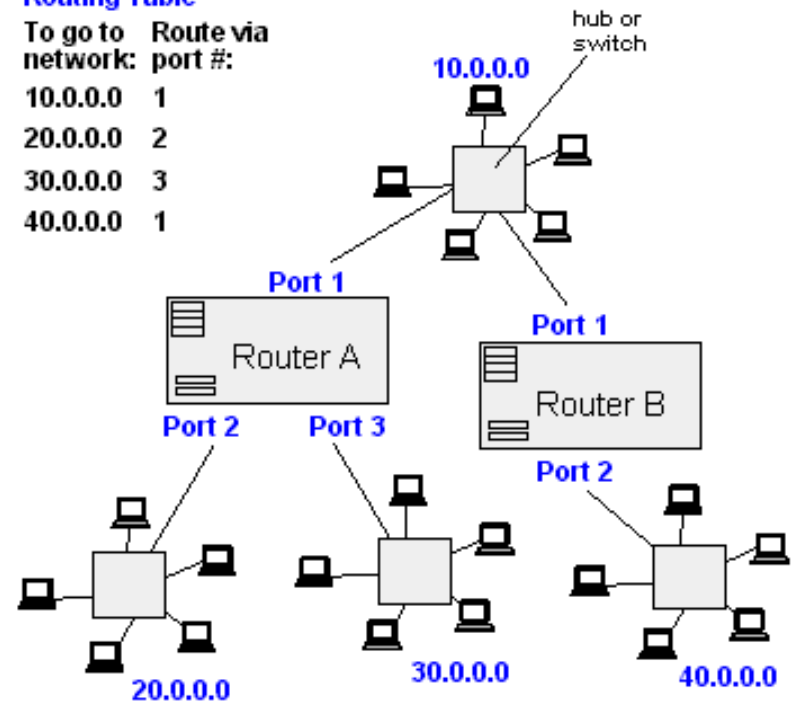
- Maps (Map Quest, Google Maps)
- Routing Systems

From Computer Desktop Encyclopedia
© 1998 The Computer Language Co. Inc.



Router A
Routing Table

To go to network:	Route via port #:
10.0.0.0	1
20.0.0.0	2
30.0.0.0	3
40.0.0.0	1



Dijkstra's algorithm

Dijkstra's algorithm - is a solution to the single-source shortest path problem in graph theory.

Works on both directed and undirected graphs. However, all edges must have nonnegative weights.

Input: Weighted graph $G=\{E,V\}$ and source vertex $v \in V$, such that all edge weights are nonnegative

Output: Lengths of shortest paths (or the shortest paths themselves) from a given source vertex $v \in V$ to all other vertices

Approach

- The algorithm computes for each vertex u the **distance** to u from the start vertex v , that is, the weight of a shortest path between v and u .
- the algorithm keeps track of the set of vertices for which the distance has been computed, called the **cloud** C
- Every vertex has a label D associated with it. For any vertex u , $D[u]$ stores an approximation of the distance between v and u . The algorithm will update a $D[u]$ value when it finds a shorter path from v to u .
- When a vertex u is added to the cloud, its label $D[u]$ is equal to the actual (final) distance between the starting vertex v and vertex u .

Dijkstra pseudocode

Dijkstra(v1, v2):

for each vertex v: // Initialization

v's distance := infinity.

v's previous := none.

v1's distance := 0.

List := {all vertices}.

while List is not empty:

v := remove List vertex with minimum distance.

mark v as known.

for each unknown neighbor n of v:

dist := v's distance + edge (v, n)'s weight.

if dist is smaller than n's distance:


n's distance := dist.

n's previous := v.

*reconstruct path from v2 back to v1,
following previous pointers.*

Example: Initialization


Distance(source) =
0



Distance (all vertices
but source) = ∞

Pick vertex in List with minimum distance.


Example: Update neighbors' distance



Distance(B) = 2


Distance(D) = 1

Example: Remove vertex with minimum distance



Pick vertex in List with minimum distance, i.e., D

Example: Update neighbors



Distance(C) = 1 + 2 = 3


Distance(E) = 1 + 2 = 3

Distance(F) = 1 + 8 = 9

Distance(G) = 1 + 4 = 5

Example: Continued...


Pick vertex in List with minimum distance (B) and update neighbors



Note : distance(D) not updated since D is already known and distance(E) not updated since it is larger than previously computed


Example: Continued...

Pick vertex List with minimum distance (E) and update neighbors




Example: Continued...

Pick vertex List with minimum distance (C) and update neighbors



Example: Continued...

Pick vertex List with minimum distance (G) and update neighbors




Previous distance




$$\text{Distance}(F) = \min(8, 5+1) = 6$$

Example (end)




Pick vertex not in S with lowest cost (F) and update neighbors

Another Example




Another Example




S: { A }


Another Example




Another Example




Another Example




Another Example




Another Example



Another Example




Another Example



Correctness : “Cloudy” Proof

When a vertex is added to the cloud, it has shortest distance to source.




- If the path to v is the next shortest path, the path to v' must be at least as long. Therefore, any path through v' to v cannot be shorter!

Minimum Spanning trees

What are trees in undirected graphs?

- We are familiar with a rooted directed tree, such as:


- Binary search tree
- Heap



- They look like:
 - You know which one the root is
 - Who the parent is
 - Who the children are


What are trees in undirected graphs?

- However, in undirected graphs, there is another definition of trees
- Tree
 - A undirected graph (V, E) , where E is the set of undirected edges
 - All vertices are connected
 - $|E| = |V| - 1$




What is a spanning tree?

- Unless otherwise specified, spanning tree is about a **tree of a undirected graph**
- A spanning tree of graph $G=(V, E)$ is:
 - A graph $G'=(V', E')$, where $V' = V, E' \subseteq E$
 - G' is a tree
- Given the following graph, are they spanning trees?



What is a spanning tree?

- Given the following graph, are they spanning trees?




Is the red sub graph a tree? Yes

Is the red sub graph a spanning tree? No. $V' \neq V$

What is a spanning tree?

- Given the following graph, are they spanning trees?




Is the red sub graph a tree? No. $|E| \neq |V'| - 1$

Is the red sub graph a spanning tree? No. It is not a tree.

What is a spanning tree?

- Given the following graph, are they spanning trees?



Is the red sub graph a tree?


No. Vertices are not all connected

Is the red sub graph a spanning tree?

No. It is not a tree.

What is a spanning tree?

- Given the following graph, are they spanning trees?




Is the red sub graph a tree? Yes

Is the red sub graph a spanning tree? Yes

What is a spanning tree?

- Given the following graph, are they spanning trees?




Is the red sub graph a tree? Yes

Is the red sub graph a spanning tree? Yes

There could be more than one spanning trees


What is a minimum spanning tree?

- When edges are weighted
- A minimum spanning tree of G is:
 - A spanning tree of G
 - With a minimum sum of a edge weights sum among all spanning trees




What is a minimum spanning tree?

- Spanning trees:



The weight sum of
this spanning tree?

48




The weight sum of
this spanning tree?

47


What is a minimum spanning tree?

- Spanning trees:



The weight sum of
this spanning tree?

37



The weight sum of
this spanning tree?

37

These 2 are minimum spanning trees
since there is no spanning tree with a
smaller total weight

Minimum spanning tree problem

- **Minimum spanning tree (MST) problem:**
given a undirected graph G , find a spanning tree with the minimum total edge weights.
- There are 2 classic greedy algorithms to solve this problem
 - Kruskal's algorithm
 - Prim's algorithm

Minimum spanning tree problem

- A application example of MST in computer network

- What if PC3 wants to communicate with PC7?

- By a routing protocol:

- PC3 broadcasts this requirement to PC2, PC4


- PC2 broadcasts it to PC1, PC5

- PC1 broadcasts it to PC5, PC6

- ...

- PC5 broadcasts it to PC2

- At least, a package is in a deadlock: PC2->PC1->PC5-
PC2->PC1->



Find and use a MST can solve this problem, since there is no cycle

Minimum spanning tree problem

- The MST algorithm can be considered as choose $|V|-1$ edges from E to form a minimum spanning tree
- Both Kruskal's algorithm and Prim's algorithm proposed a greedy choice about how to choose a edge
- They have greedy choice property because of textbook **Corollary 23.2**


Kruskal's algorithm

- Input: $G=(V, E)$
- Output: a minimum spanning tree (V, A)
- **Kruskal** ($G=(V, E)$)
 - Set each vertices in V as a vertex set;
 - Set A as an empty set;
 - while (there are more than 1 vertex set){
 - Add a smallest edge (u, v) to A where u, v are not in the same set;
 - Merge the set contains u with the set contains v ;
 - }
- In most of the implementations, edges are sorted on weights at the beginning, and then scan them once
- The running time = $O(|E| \lg |V|)$



Kruskal's algorithm

Kruskal ($G=(V, E)$)

Set each vertices in V as a vertex set;
 Set A as an empty set;
 while (there are more than 1 vertex set){
 Add a smallest edge (u, v) to A where u ,
 v are not in the same set;
 Merge the set contains u with the set
 contains v ;
 }



Total weight: 37

Edge	Weight
 (h, g) 	1
 (i, c) 	2
 (g, f) 	2
 (c, f) 	4
 (a, b) 	4
 (i, g)	6
 (h, i)	7
 (c, d) 	7
 (a, h) 	8
 (b, c)	8
 (d, e) 	9
(e, f)	10
(b, h)	11
(d, f)	14

Prim's algorithm

- Input: $G=(V, E)$
- Output: a minimum spanning tree (V, A)
- **Prim**($G=(V, E)$)
 - Set A as an empty set;
 - Choose a vertex as the current MST;
 - while (not all the vertices are added into the MST){
 - Choose the smallest edge crossing current MST vertices and other vertices, and put it to MST ;
 - }
- The running time = $O(|E| \lg |V|)$

Prim's algorithm

Prim($G=(V, E)$)


Set A as an empty set;

Choose a vertex as the current MST;

while (not all the vertices are added into the MST){

 Choose the smallest safe edge, and put it to MST;

}



Total weight: 37

Summary

- Prim's algorithm and Kruskal's algorithm have the same asymptotically running time